

Dheeren Tejani

Email: dheerenntejani@gmail.com | Phone: +91-7822868074

GitHub: <https://github.com/dheeren-tejani> | LinkedIn: <https://www.linkedin.com/in/dheeren-tejani>

Portfolio: <https://dheerentejani.netlify.app> | Location: Mumbai, India

Summary

Final-year AI student specializing in Computer Vision and ML System Optimization. Built custom super-resolution architectures achieving $7\times$ inference speedup through TensorRT optimization, fine-tuned Stable Diffusion models using PEFT/LoRA on low-VRAM hardware, and deployed full-stack applications (FastAPI + React) for real-time AI inference. Seeking ML Engineer/Computer Vision roles to apply production optimization and deployment expertise.

Education

B.Sc. in Artificial Intelligence

Jun 2023 – Apr 2026 (Expected)

College : Vivekanand Education Society's College of Arts, Science and Commerce (Autonomous) Mumbai

CGPA : 8.65

Skills

Languages: Python, JavaScript/TypeScript, SQL

AI/ML: PyTorch, TensorRT, Computer Vision, NLP, Super-Resolution, Model Optimization

Deployment: FastAPI, Streamlit, Docker, Git, AWS

Web Dev: React, TypeScript, Tailwind CSS

Data/Tools: NumPy, Pandas, OpenCV, Matplotlib

Projects

AI Super-Resolution System (Image + Video)

PyTorch, TensorRT, FastAPI, React

- Engineered an end-to-end video enhancement suite that achieved a **$7\times$ inference speedup** on an RTX 3050 (4GB VRAM) by optimizing a CNN/GAN model with TensorRT, multithreading, and GPU offloading techniques.
- Developed and deployed a full-stack application with a React frontend and FastAPI backend for real-time, low-latency image and video upscaling.

Diabetic Retinopathy Detection

Python, Computer Vision, Streamlit

- Fine-tuned a ResNet50 classifier on the APTOS 2019 dataset, achieving **near 75% accuracy and a Quadratic Weighted Kappa of 0.80** for retinal image severity classification.
- Integrated **Grad-CAM overlays** and interactive severity charts into a Streamlit dashboard to improve model interpretability for potential clinical use.

Cyberpunk Stable Diffusion Fine-Tune

PEFT (LoRA), bitsandbytes

- Fine-tuned the Stable Diffusion 1.5 model by applying Low-Rank Adaptation (LoRA) on a custom-curated dataset to generate high-quality, **cyberpunk-style images**.
- Implemented parameter-efficient training (PEFT) techniques with ‘bitsandbytes’ to enable fine-tuning on consumer-grade, low-VRAM GPU (16 GB VRAM T4 GPU).

Other Activity (Hackathon)

Odoo Hackathon: Led the development of a StackOverflow-style Q&A prototype, implementing Google OAuth for user authentication and building the core application logic using Python and JavaScript